

ISYE 6420 - BUGS to Stan

Josh Blakely

2/15/23

Table of contents

Introduction	5
Why?	6
Site Structure	7
Plans	8
 I UNIT 1	 9
About this course and website	10
Other recommended resources	10
Textbooks and courses	10
Blogs	11
Podcasts	11
Other	11
 Topics Covered	 12
Homework 1	12
Homework 2	12
Homework 3	12
Homework 4	13
Midterm	13
Homework 5	13
Homework 6	13
Final	13
Project	13
 II UNIT 2	 14
Placeholder	15

III UNIT 3	16
Placeholder	17
IV UNIT 4	18
Placeholder	19
V UNIT 5	20
Placeholder	21
VI UNIT 6	22
Stress, Diet, and Plasma Acids	23
Problem Statement	23
BUGS step function	24
How do I track non-random variables in Stan?	24
Stress, Diet, and Plasma Acids Model	24
Sampling	25
Dugongs	27
Problem Statement	27
References	27
Input Data	27
Sampling	29
Equivalence of Generic and Brand-name Drugs	31
Problem Statement	31
Input Data	31
How do I track non-random variables in Stan?	32
Stress, Diet, and Plasma Acids Model	32
Sampling	32
Psoriasis: Two Sample Problem - paired data	34
Problem Statement	34
Model	35
Sampling	35
Taste of Cheese	37
Problem Statement	37
Model	37

Sampling	39
Predictions	39
VII UNIT 7	41
Priors as Hidden Mixtures	42
Problem Statement	42
Model	42
Sampling	43
Coagulation	45
Problem Statement	45
Model	45
Sampling	47
Results	47
Plot Intervals	48
VIIIUNIT 8	50
Placeholder	51
IX UNIT 9	52
Placeholder	53
X UNIT 10	54
Placeholder	55

Introduction

This repository contains R translations of examples from Georgia Tech's ISYE 6420: Bayesian Statistics, created by Professor Brani Vidakovic and currently taught by Professor Roshan Joseph and head TA Greg Schreiter. It also has additional notes on each lecture.

Why?

When I was taking this course the only option for Mac users was to either use a virtual machine to run WinBUGS/OpenBUGS. During the semester I was taking this, Aaron Reding was creating an alternate using Python and PyMC (see [here](#)), but I found that PyMC had too many extra steps and was still fairly buggy.

After quite a bit of research, I stumbled across Stan. The benefits that Stan provides over PyMC is that it can be run using Python ([PyStan](#) or `CmdStanPy`) or R ([CmdStanR](#) or [RStan](#)) or MATLAB, most of the models would also look similar to WinBUGS/OpenBUGS and finally the documentation for [Stan](#) is excellent.

When I became a TA for this course I noticed that not a lot of people were using Stan, so I wanted to give the students an alternative to running Markov Chain Monte Carlo for Bayesian statistics.

Site Structure

This site was built with [Quarto Book](#) and is made up of a combination of Stan files and Quarto Document (which is similar to RMarkdown files). You can view the entire repository on [GitHub](#).

All pages match a corresponding lecture on Canvas, except when there are more pages than lectures in a unit. In which case the additional pages will be at the end of the unit. For example there are only eight lectures in Unit 3 on Canvas. The first eight pages here match the lecture numbers, while the ninth page has a supplementary problem and wasn't part of the original lectures.

Any necessary data will either have a download link or, if the data is compact enough, will be included in the code.

Plans

As of Summer 2023, this is still a work in progress. I have been working through each unit trying to add the important information. Stan is heavily used from Unit 6 on, so that has been where my focus has been aligned.

Part I

UNIT 1

About this course and website

The course was originally developed at Georgia Tech in 2004 by Professor Brani Vidakovic, who is now the head of the [Texas A&M Department of Statistics](#). Many of the individual examples are related to biostatistics (Prof. Vidakovic wrote the textbook *Engineering Biostatistics*), but the methods are broadly applicable.

See also Professor Vidakovic's [publication history](#).

[Professor Joseph's](#).

This site mostly follows the original [course outline](#). Each example lists the corresponding lecture video and contains a download link to the original code file(s). Only the lecture video where the professor goes over the example code will be listed, but there may be other relevant lectures that you'll need to watch. Any necessary data will either have a download link or, if the data is compact enough, will be included in the code.

This is a supplement to the Canvas site, not a replacement!

Other recommended resources

These are not required for the class, but they might be helpful. Prof. Vidakovic's lectures often assume that the student has a certain amount of background knowledge, so if you feel lost or if you just want to dive deeper into the subject check them out.

Textbooks and courses

Statistical Rethinking by Richard McElreath is a great book for gaining intuition about Bayesian inference and modeling in general.

- [main site](#)
- [lecture videos](#)
- [R and Stan code examples](#).

Bayesian Data Analysis by Gelman, Carlin, Stern, Dunson, Vehtari, and Rubin goes into more mathematical theory than *Statistical Rethinking*. I use it as a reference - not planning to try reading this one all the way through!

- [main site](#)
- [pdf](#)

Blogs

- Andrew Gelman, author of *Bayesian Data Analysis* (above), has a blog: [Statistical Modeling, Causal Inference, and Social Science](#).
- [Dan Simpson](#), one of the maintainers of the [Stan PPL](#), has a blog called [Un garçon pas comme les autres \(Bayes\)](#) with opinionated and funny deep dives into various Bayesian topics. Warning: NSFW language.
- [Count Bayesie: Probably a probability blog](#) by Will Kurt.
- Michael Betancourt, another Stan developer, has a [series of incredibly in-depth](#) posts and notebooks on Bayesian modeling.
- PyMC developer [Austin Rochford's blog](#) has a lot of good posts.
- Another PyMC developer, [Oriol Abril](#), posts some really helpful PyMC examples.

Podcasts

- [Learn Bayes Stats](#) by Alex Andorra, one of the PyMC developers.

Other

- [Stan User's Guide](#)
- [Aaron Reding's PyMC](#)

Topics Covered

This lecture goes over the [topics covered](#) in the course. Here are some notes about each homework and the corresponding units.

Homework 1

Units covered: 1–3

We start out with a very brief review of probability and an introduction to Bayes' rule. The first homework generally has questions that don't require any calculus, just basic probability rules. They may include circuit problems, use of Bayes' Rule, the Law of Total Probability, sensitivity and specificity, positive or negative predictive value (PPV/NPV), and simple Bayesian networks. These problems can be trickier than they seem at first!

Homework 2

Lessons covered: 4.1–4.9

We start to discuss random variables, distributions, and Bayesian inference with conjugate cases.

Homework 3

Lessons covered: 4.10–4.19

Here we get into hypothesis testing, ways to describe the posterior, different priors and their effects, and a brief discussion of Empirical Bayes methods. We're still using conjugate cases to find posteriors.

Homework 4

Unit covered: 5

This unit we mostly move away from conjugate problems and start discussing MCMC methods for sampling from the posterior. Usually we'll have one Metropolis-Hastings and one Gibbs sampler question.

Midterm

Units covered: 1–5

The midterm could include anything from the first half, but tends to focus on Units 4 and 5.

Homework 5

Units covered: 6–7

This is where we start using Probabilistic Programming Languages (PPLs) like BUGS, [Stan](#), or [PyMC](#). We start with linear regression and move into hierarchical models.

Homework 6

Units covered: 8–9

Unit 8 discusses missing and censored data along with time-to-event models. Unit 9 goes into model building, selection, and checking.

Final

Units covered: 1–10

The final is comprehensive but tends to focus on concepts from the second half of the course.

Project

There's also an open-ended project. It could potentially be based on any part of the course, but most students opt to try out or extend the models from the second half of the course.

Part II

UNIT 2

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.

Part III

UNIT 3

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.

Part IV

UNIT 4

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.

Part V

UNIT 5

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.

Part VI

UNIT 6

Stress, Diet, and Plasma Acids

```
library(cmdstanr)
```

This example introduces tracking of deterministic variables and shows how to recreate the BUGS step function in Stan.

It is adapted from [Unit 6: stressacids.odc](#).

Associated lecture video: Unit 6 lesson 5

Problem Statement

In the study [Interrelationships Between Stress, Dietary Intake, and Plasma Ascorbic Acid During Pregnancy](#) conducted at the Virginia Polytechnic Institute and State University, the plasma ascorbic acid levels of pregnant women were compared for smokers versus non-smokers. Thirty-two women in the last three months of pregnancy, free of major health disorders, and ranging in age from 15 to 32 years were selected for the study. Prior to the collection of 20 ml of blood, the participants were told to avoid breakfast, forego their vitamin supplements, and avoid foods high in ascorbic acid content. From the blood samples, the plasma ascorbic acid values of each subject were determined in milligrams per 100 milliliters.

I start with the data pasted from `stressacids.odc`, then create one list for smokers and one for nonsmokers.

```
plasma <- c(0.97, 0.72, 1.00, 0.81, 0.62, 1.32, 1.24, 0.99, 0.90, 0.74,  
            0.88, 0.94, 1.06, 0.86, 0.85, 0.58, 0.57, 0.64, 0.98, 1.09,  
            0.92, 0.78, 1.24, 1.18, 0.48, 0.71, 0.98, 0.68, 1.18, 1.36,  
            0.78, 1.64)  
  
smo <- c(1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
         1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2)
```

```
nonsmoker_index <- which(smo == 1)
plasma_smokers <- plasma[-nonsmoker_index]
plasma_nonsmokers <- plasma[nonsmoker_index]
```

BUGS step function

BUGS defines the step function like this:

$$step(e) = \begin{cases} 1, & e \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

Stan follows what you would expect in a programming language. We implement this like:

```
e >= 0 ? 1 : 0;
```

How do I track non-random variables in Stan?

One nice thing about BUGS is you can easily track both deterministic and non-deterministic variables while sampling. For Stan, you add the variable to the **generated quantities** block.

```
mod <- cmdstan_model(stress_diet_stan)
```

Stress, Diet, and Plasma Acids Model

```
data {
  int<lower=0> N_smoker;
  int<lower=0> N_nonsmoker;
  vector[N_smoker] plasma_smoker;
  vector[N_nonsmoker] plasma_nonsmoker;
}

parameters {
  real<lower=0> tau_nonsmoker;
  real mu_nonsmoker;
  real<lower=0> tau_smoker;
  real mu_smoker;
}

model {
```



```

tau_nonsmoker ~ gamma(0.0001, 0.0001);
tau_smoker ~ gamma(0.0001, 0.0001);
mu_nonsmoker ~ normal(0, 100); // equivalent to BUGS tau = 0.0001
mu_smoker ~ normal(0, 100);

plasma_smoker ~ normal(mu_smoker, 1 / sqrt(tau_smoker));
plasma_nonsmoker ~ normal(mu_nonsmoker, 1 / sqrt(tau_nonsmoker));
}

generated quantities {
  real testmu;
  real r;

  testmu = (mu_smoker >= mu_nonsmoker ? 1 : 0);
  r = tau_nonsmoker / tau_smoker;
}

```

Sampling

Let us prepare the data to be passed into sample

```

input_data <- list(N_smoker=length(plasma_smokers),
                  N_nonsmoker=length(nonsmoker_index),
                  plasma_smoker=plasma_smokers,
                  plasma_nonsmoker=plasma_nonsmokers
)

```

Now that we have the data to pass into our sampling, let's proceed.

```

fit <- mod$sample(
  data = input_data,
  seed = 123,
  chains = 4,
  parallel_chains = 4,
  refresh = 500, # print update every 500 iterations
  iter_warmup = 1000,
  iter_sampling = 5000
)

fit$summary()

```

```
# A tibble: 7 x 10
  variable      mean median      sd    mad     q5    q95  rhat ess_bulk ess_tail
  <chr>      <num>  <num>  <num>  <num>  <num>  <num> <num>    <num>    <num>
1 lp__        27.8   28.1   1.52   1.31   24.9   29.6   1.00    7997.    9743.
2 tau_nonsmok~ 22.6   21.9   6.68   6.53   12.9   34.5   1.00   17854.   12231.
3 mu_nonsmoker 0.912   0.912  0.0449 0.0436  0.839   0.986   1.00   17455.   12780.
4 tau_smoker   6.53    5.92   3.50   3.20    2.03   13.1   1.00   13257.   10335.
5 mu_smoker    0.977   0.977  0.162   0.145   0.721   1.24   1.00   13512.    9947.
6 testmu       0.667    1     0.471    0       0       1     1.00   16274.     NA
7 r            4.83    3.72   4.27    2.19    1.43   11.7   1.00   14304.   11261.
```

These results are very similar to PyMC results.

Dugongs

```
library(cmdstanr)
```

This example is the first example of dealing with missing data.

It is adapted from [Unit 6: dugongsmissing.odc](#).

Associated lecture video: Unit 6 lesson 6

Problem Statement

Carlin and Gelfand (1991) investigated the age (x) and length (y) of 27 captured dugongs (sea cows). Estimate parameters in a nonlinear growth model.

References

Data provided by Ratkowsky (1983).

Carlin, B. and Gelfand, B. (1991). An Iterative Monte Carlo Method for Nonconjugate Bayesian Analysis, *Statistics and Computing*,

Ratkowsky, D. (1983). Nonlinear regression modeling: A unified practical approach. M. Dekker, NY, viii, 276 p.

Input Data

```
X <- c(1.0, 1.5, 1.5, 1.5, 2.5, 4.0, 5.0, 5.0, 7.0, 8.0, 8.5, 9.0, 9.5,  
      9.5, 10.0, 12.0, 12.0, 13.0, 13.0, 14.5, 15.5, 15.5, 16.5, 17.0,  
      22.5, 29.0, 31.5)  
y <- c(1.80, 1.85, 1.87, -1, 2.02, 2.27, 2.15, 2.26, 2.47, 2.19, 2.26,  
      2.40, 2.39, 2.41, 2.50, 2.32, 2.32, 2.43, 2.47, 2.56, 2.65, 2.47,  
      2.64, 2.56, 2.70, 2.72, -1)
```

Stan imputes missing values differently from PyMC and Bugs. We have to pass in the indices for the missing values as well as the count of observed and missing values.

```
mod <- cmdstan_model(dugongs_stan)
```

now that we have compiled our stan file, we can print out the model that we will use for this:

```
mod$print()
```

```
data {
  int<lower=0> N_obs;
  int<lower=0> N_mis;
  array[N_obs] int<lower=1, upper=N_obs + N_mis> iy_obs; // index of obs y's
  array[N_mis] int<lower=1, upper=N_obs + N_mis> iy_mis; // index of mis y's
  vector[N_obs + N_mis] X;
  vector[N_obs] y_obs; // actual y values that were observed
}

transformed data {
  int<lower=0> N = N_obs + N_mis; // total size of the input
}

parameters {
  real alpha;
  real beta;
  real gamma;
  real sigma;
  vector[N_mis] y_mis; // account for missing y
}

transformed parameters {
  vector[N] Y; // Imputed Y values
  Y[iy_obs] = y_obs; // actual y values
  Y[iy_mis] = y_mis; // missing y values
}

model {
  // priors
  alpha ~ uniform(0, 100);
  beta ~ uniform(0, 100);
  gamma ~ uniform(0, 1);
  sigma ~ uniform(-10, 10);
}
```

```
Y ~ normal(alpha - beta * pow(gamma, X), sigma);
}
```

Sampling

Let us prepare the data to be passed into sample

```
# Initial data has `-1` where missing.
idx.mis <- which(y == -1)
idx.obs <- which(y != -1)

input_data <- list(N_obs=length(idx.obs), N_mis=length(idx.mis),
                  iy_obs=idx.obs, iy_mis=idx.mis,
                  X=X,
                  y_obs=y[idx.obs]
                )
```

Now that we have the data to pass into our sampling, let's proceed.

```
fit <- mod$sample(
  data = input_data,
  seed = 123,
  chains = 4,
  parallel_chains = 4,
  refresh = 500, # print update every 500 iterations
  iter_warmup = 1000,
  iter_sampling = 5000
)

fit$summary()
```

A tibble: 34 x 10

	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk	ess_tail
	<chr>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>
1	lp__	49.5	49.9	2.09	1.90	45.5	52.2	1.00	5446.	7258.
2	alpha	2.73	2.71	0.125	0.106	2.57	2.96	1.00	3806.	2567.
3	beta	0.986	0.978	0.105	0.0925	0.833	1.16	1.00	5014.	2863.
4	gamma	0.886	0.891	0.0358	0.0313	0.823	0.936	1.00	4033.	3205.
5	sigma	0.100	0.0980	0.0165	0.0153	0.0771	0.130	1.00	8360.	8205.
6	y_mis[1]	1.91	1.91	0.114	0.111	1.72	2.09	1.00	11392.	11007.

7	y_mis[2]	2.69	2.69	0.128	0.123	2.48	2.90	1.00	6343.	6816.
8	Y[1]	1.8	1.8	0	0	1.8	1.8	NA	NA	NA
9	Y[2]	1.85	1.85	0	0	1.85	1.85	NA	NA	NA
10	Y[3]	1.87	1.87	0	0	1.87	1.87	NA	NA	NA

i 24 more rows

These results are very similar to PyMC results.

Equivalence of Generic and Brand-name Drugs

```
library(cmdstanr)
```

This example introduces tracking of deterministic variables and shows how to recreate the BUGS step function in Stan.

It is adapted from [Unit 6: equivalence.odc](#).

Associated lecture video: Unit 6 lesson 7

Problem Statement

The manufacturer wishes to demonstrate that their generic drug for a particular metabolic disorder is equivalent to a brand name drug. One of indication of the disorder is an abnormally low concentration of levocarnitine, an amino acid derivative, in the plasma. The treatment with the brand name drug substantially increases this concentration.

A small clinical trial is conducted with 43 patients, 18 in the Brand Name Drug arm and 25 in the Generic Drug arm. The increases in the log-concentration of levocarnitine are in the data below.

The FDA declares that bioequivalence among the two drugs can be established if the difference in response to the two drugs is within 2 units of log-concentration. Assuming that the log-concentration measurements follow normal distributions with equal population variance, can these two drugs be declared bioequivalent within a tolerance ± 2 units?

Input Data

Starting with the data pasted from `equivalence.odc`.

```
increase.1 <- c(7, 8, 4, 6, 10, 10, 5, 7, 9, 8, 6, 7, 8, 4, 6, 10, 8, 9)
increase.2 <- c(6, 7, 5, 9, 5, 5, 3, 7, 5, 10, 8, 5, 8, 4, 4, 8, 6, 11, 7, 5, 5, 5, 7, 4,
```

How do I track non-random variables in Stan?

One nice thing about BUGS is you can easily track both deterministic and non-deterministic variables while sampling. For Stan, you add the variable to the **generated quantities** block.

```
mod <- cmdstan_model(equivalence_stan)
```

Stress, Diet, and Plasma Acids Model

```
data {
  int<lower=0> N;
  int<lower=0> M;
  vector[N] y_type1;
  vector[M] y_type2;
}

parameters {
  real mu_type1;
  real mu_type2;
  real prec;
}

transformed parameters {
  real mudiff = mu_type1 - mu_type2;
  real sigma = 1 / sqrt(prec);
  real probint = ((mudiff + 2) >= 0 ? 1 : 0) * ((2 - mudiff) >= 0 ? 1 : 0);
}

model {
  mu_type1 ~ normal(10, 1 / sqrt(1e-5));
  mu_type2 ~ normal(10, 1 / sqrt(1e-5));
  prec ~ gamma(0.001, 0.001);

  y_type1 ~ normal(mu_type1, sigma);
  y_type2 ~ normal(mu_type2, sigma);
}
```

Sampling

Let us prepare the data to be passed into sample


```

input_data <- list(N=length(increase.1),
                  M=length(increase.2),
                  y_type1=increase.1,
                  y_type2=increase.2
)

```

Now that we have the data to pass into our sampling, let's proceed.

```

fit <- mod$sample(
  data = input_data,
  seed = 123,
  chains = 4,
  parallel_chains = 4,
  refresh = 500, # print update every 500 iterations
  iter_warmup = 1000,
  iter_sampling = 5000
)

```

```
fit$summary()
```

```
# A tibble: 7 x 10
```

	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk	ess_tail
	<chr>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>
1	lp__	-49.4	-49.1	1.22	0.978	-51.8	-48.1	1.00	9921.	12770.
2	mu_type1	7.33	7.34	0.467	0.461	6.57	8.09	1.00	16880.	13439.
3	mu_type2	6.20	6.21	0.400	0.393	5.55	6.86	1.00	16065.	13287.
4	prec	0.263	0.259	0.0581	0.0575	0.176	0.365	1.00	15209.	11437.
5	mudiff	1.13	1.13	0.617	0.616	0.122	2.14	1.00	16797.	13707.
6	sigma	1.99	1.96	0.226	0.218	1.66	2.39	1.00	15209.	11437.
7	probint	0.922	1	0.268	0	0	1	1.00	14222.	NA

These results are very similar to PyMC results and BUGS results.

Psoriasis: Two Sample Problem - paired data

```
library(cmdstanr)
```

This is our first example of hypothesis testing.

It is adapted from [Unit 6: psoriasis.odc](#).

Associated lecture video: Unit 6 lesson 7

Problem Statement

Woo and McKenna (2003) investigated the effect of broadband ultraviolet B (UVB) therapy and topical calcipotriol cream used together on areas of psoriasis. One of the outcome variables is the Psoriasis Area and Severity Index (PASI), where a lower score is better.

The PASI scores for 20 subjects are measured at baseline and after 8 treatments. Do these data provide sufficient evidence to indicate that the combination therapy reduces PASI scores?

Classical Analysis:

```
d = baseline - after;
n = length(d);
dbar = mean(d);    dbar = 6.3550
sdd = sqrt(var(d)); sdd = 4.9309
tstat = dbar / (sdd / sqrt(n)); tstat = 5.7637
```

Reject H_0 at the level $\alpha = 0.05$ since the $p_value = 0.00000744 < 0.05$

95% CI is [4.0472, 8.6628]

See [Unit 6: Stress, Diet and Plasma Acids](#) to find out more about recreating the BUGS step function.

```
baseline <- c(5.9, 7.6, 12.8, 16.5, 6.1, 14.4, 6.6, 5.4, 9.6, 11.6, 11.1, 15.6, 9.6, 15.2,
after <- c(5.2, 12.2, 4.6, 4, 0.4, 3.8, 1.2, 3.1, 3.5, 4.9, 11.1, 8.4, 5.8, 5, 6.4, 0, 2.
```

```
mod <- cmdstan_model(psoriasis_stan)
```

We do get a decent amount of warnings, but the model compiles and runs.

Model

```
// Psoriasis: Two Sample Problem - Paired Data
data {
  int<lower=0> N;
  vector[N] baseline;
  vector[N] after;
}

transformed data {
  vector[N] diff = baseline - after;
}

parameters {
  real mu;
  real prec;
}

transformed parameters {
  real sigma = 1 / sqrt(prec);
  real ph1;
  ph1 = (mu >= 0 ? 1 : 0);
}

model {
  mu ~ normal(0, 316);
  prec ~ gamma(0.001, 0.001);

  diff ~ normal(mu, sigma);
}
```

Sampling

Let us prepare the data to be passed into sample

```
input_data <- list(N=length(baseline),
                  baseline=baseline,
                  after=after
)
```

Now that we have the data to pass into our sampling, let's proceed.

```
fit <- mod$sample(
  data = input_data,
  seed = 123,
  chains = 4,
  parallel_chains = 4,
  refresh = 500, # print update every 500 iterations
  iter_warmup = 1000,
  iter_sampling = 5000
)
```

```
fit$summary()
```

```
# A tibble: 5 x 10
  variable    mean  median    sd    mad    q5    q95  rhat ess_bulk
  <chr>      <num>   <num> <num> <num> <num> <num> <num>   <num>
1 lp__    -39.2   -38.9  1.02  0.714 -41.2  -38.3    1.00   8495.
2 mu       6.37    6.37  1.18  1.12   4.45   8.27    1.00   9440.
3 prec     0.0411  0.0398 0.0133 0.0129 0.0218 0.0651    1.00   9661.
4 sigma     5.14    5.01  0.893  0.814   3.92   6.78    1.00   9661.
5 phi       1        1      0      0        1        1      NA      NA
# i 1 more variable: ess_tail <num>
```

These results are very similar to BUGS and [PyMC results](#).

Taste of Cheese

```
library(cmdstanr)
```

Adapted from Unit 6:

The data was downloaded from [here](#) and there is a copy [here](#). If you wish to download the data, then right click and save the link as a csv file.

Problem Statement

As cheddar cheese matures, a variety of chemical processes take place. The taste of matured cheese is related to the concentration of several chemicals in the final product. In a study of cheddar cheese from the LaTrobe Valley of Victoria, Australia, samples of cheese were analyzed for their chemical composition and were subjected to taste tests. Overall taste scores were obtained by combining the scores from several tasters.

Can the score be predicted well by the predictors: Acetic, H2S, and Lactic?

```
df <- read.csv(file.path '..', 'data', 'cheese.csv'), header = T, colClasses = c("NULL", N

# This data list will be passed into the stan model
data_list <- list(
  N = dim(df)[1],
  acetic = df$Acetic,
  h2s = df$H2S,
  lactic = df$Lactic,
  y = df$taste
)
```

Model

We will compile/build the stan model by running `cmdstan_model`. The `cheese_program` is the file path for the stan file, i.e. `../cheese.stan`.

```
mod <- cmdstan_model(cheese_stan)
```

Now that we have compiled our stan file, we can print out the model:

```
mod$print()
```

```
// The data that is input into the model
data {
  int<lower=0> N;
  vector[N] acetic;
  vector[N] h2s;
  vector[N] lactic;
  vector[N] y;
}

parameters {
  real b0;           // Intercept coefficient
  real b1;           // Slope coefficient
  real b2;
  real b3;
  real tau;
}

model {
  b0 ~ normal(0, 1 / sqrt(1e-5));
  b1 ~ normal(0, 1 / sqrt(1e-5));
  b2 ~ normal(0, 1 / sqrt(1e-5));
  b3 ~ normal(0, 1 / sqrt(1e-5));
  tau ~ gamma(0.001, 0.001);

  y ~ normal(b0 + b1 * acetic + b2 * h2s + b3 * lactic, 1 / sqrt(tau));
}

generated quantities {
  array[N] real y_hat;
  y_hat = normal_rng(b0 + b1 * acetic + b2 * h2s + b3 * lactic, 1 / sqrt(tau));
}
```

Sampling

Now that we have created the model above and compiled the stan program, we can start sampling.

```
fit <- mod$sample(  
  data = data_list,  
  seed = 1,  
  chains = 4,  
  parallel_chains = 4,  
  refresh = 500, # print update every 500 iterations  
  iter_warmup = 1000,  
  iter_sampling = 5000 # iterate 5,000 times  
)  
  
fit$summary(variables = c('b0', 'b1', 'b2', 'b3', 'tau'))
```

```
# A tibble: 5 x 10  
  variable      mean    median      sd      mad      q5      q95    rhat ess_bulk  
  <chr>      <num>    <num>    <num>    <num>    <num>    <num> <num>    <num>  
1 b0      -28.6    -28.6    20.6    20.3    -6.26e+1  4.76    1.00    7508.  
2 b1       0.262     0.232     4.68     4.59    -7.32e+0  7.89    1.00    7056.  
3 b2       3.92      3.91      1.30      1.29     1.78e+0  6.04    1.00    8871.  
4 b3      19.7      19.6      9.02      8.70     4.87e+0 34.4    1.00    9316.  
5 tau      0.00973    0.00949    0.00270    0.00267  5.75e-3  0.0145    1.00    11940.  
# i 1 more variable: ess_tail <num>
```

The results are pretty close to OpenBUGS.

Predictions

In order to get the predictions, we will just need to pass in a new list of parameters with the predicted value

```
pred_data <- list(  
  N = 1,  
  acetic = 5.0,  
  h2s = 7.1,  
  lactic = 1.5,  
  y = 0 # We don't care what the value is for this because we are just trying to predict
```

```
)  
  
pred <- mod$generate_quantities(fitted_params = fit, data = pred_data)  
  
print(pred)
```

```
variable  mean median    sd   mad    q5   q95  
y_hat[1] 30.03  29.99 11.11 10.67 11.92 48.22
```


Part VII

UNIT 7

Priors as Hidden Mixtures

```
library(cmdstanr)
```

It is adapted from [Unit 7: tasmixture.odc](#).

Problem Statement

We are comparing the Student T likelihood with Normal prior vs Normal likelihood with Gamma prior.

Model

```
mod <- cmdstan_model(tas_mixture)
```

now that we have compiled our stan file, we can print out the model that we will use for this:

```
mod$print()
```

```
data {  
  real df;  
  real mu0;  
  real tau1;  
  real tau;  
  real X;  
  real Y;  
}
```

```
parameters {  
  real mu1;  
  real mu2;  
  real prec;  
}
```

```

transformed parameters {
  real alpha = df / 2;
  real beta = df / (2.0 * tau);
}

model {
  mu1 ~ student_t(df, mu0, tau);
  X ~ normal(mu1, 1 / sqrt(tau1));

  prec ~ gamma(alpha, beta);
  mu2 ~ normal(mu0, 1 / sqrt(prec));
  Y ~ normal(mu2, 1 / sqrt(tau1));
}

```

Sampling

We don't technically need to send anything in. We could add the values to the [transformed data](#) block. For this I decided to pass in the value. Therefore, let us prepare the data to be passed into sample

```

input_data <- list(
  df = 6,
  mu0 = 6,
  tau1 = 10,
  tau = 0.4,
  X = 10,
  Y = 10
)

```

Now that we have the data to pass into our sampling, let's proceed.

```

fit <- mod$sample(
  data = input_data,
  seed = 123,
  chains = 4,
  parallel_chains = 4,
  refresh = 500, # print update every 500 iterations
  iter_warmup = 1000,
  iter_sampling = 5000
)

```

```
fit$summary()
```

```
# A tibble: 6 x 10
```

	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk	ess_tail
	<chr>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>
1	lp__	-13.1	-12.8	1.26	1.02	-15.6	-11.8	1.00	7497.	9707.
2	mu1	9.83	9.83	0.323	0.321	9.29	10.4	1.00	13632.	11613.
3	mu2	9.91	9.91	0.315	0.309	9.40	10.4	1.00	13427.	11235.
4	prec	0.234	0.211	0.126	0.115	0.0701	0.475	1.00	8902.	6831.
5	alpha	3	3	0	0	3	3	NA	NA	NA
6	beta	7.5	7.5	0	0	7.5	7.5	NA	NA	NA

These results are very similar to PyMC results.

Coagulation

```
library(cmdstanr)
library(bayesplot)
```

An example of Bayesian ANOVA

It is adapted from [Unit 7: anovacoagulation.odc](#).

Problem Statement

Here 24 animals are randomly allocated to 4 different diets, but the numbers allocated to different diets are not the same. The coagulation time for blood is measured for each animal. Are the diet-based differences significant?

[Box, Hunter, Hunter; Statistics for Experimenters, p. 166](#)

Model

```
mod <- cmdstan_model(anovacoagulation)
```

As you can see, this isn't the best implementation of this. It currently requires typing each variable multiple times. This could have loops, which would reduce having to type out each variable multiple times.

```
mod$print()
```

```
data {
  int<lower=0> N1;
  int<lower=0> N2;
  int<lower=0> N3;
  int<lower=0> N4;
  vector[N1] y1;
  vector[N2] y2;
```

```

    vector[N3] y3;
    vector[N4] y4;
}

parameters {
    real mu0;
    real<lower=0> tau;

    real alpha4;
    real alpha3;
    real alpha2;
}

transformed parameters {
    real alpha1 = -(alpha2 + alpha3 + alpha4);
    real mu1 = mu0 + alpha1;
    real mu2 = mu0 + alpha2;
    real mu3 = mu0 + alpha3;
    real mu4 = mu0 + alpha4;
}

model {
    mu0 ~ normal(0, 1 / sqrt(0.0001));
    tau ~ gamma(0.001, 0.001);

    alpha4 ~ normal(0, 1 / sqrt(0.0001));
    alpha3 ~ normal(0, 1 / sqrt(0.0001));
    alpha2 ~ normal(0, 1 / sqrt(0.0001));

    y1 ~ normal(mu1, 1 / sqrt(tau));
    y2 ~ normal(mu2, 1 / sqrt(tau));
    y3 ~ normal(mu3, 1 / sqrt(tau));
    y4 ~ normal(mu4, 1 / sqrt(tau));
}

generated quantities {
    real onetwo = alpha1 - alpha2;
    real onethree = alpha1 - alpha3;
    real onefour = alpha1 - alpha4;
    real twothree = alpha2 - alpha3;
    real twofour = alpha2 - alpha4;
    real threefour = alpha3 - alpha4;
}

```

Sampling

```
# cut and pasted data from .odc file
times = c(62, 60, 63, 59, 63, 67, 71, 64, 65, 66, 68, 66, 71, 67, 68,
          68, 56, 62, 60, 61, 63, 64, 63, 59)
diets = c(1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 3, 3, 3, 3, 3,
          3, 4, 4, 4, 4, 4, 4, 4, 4)

y1 <- times[diets == 1]
y2 <- times[diets == 2]
y3 <- times[diets == 3]
y4 <- times[diets == 4]

input_data <- list(
  N1 = length(y1),
  N2 = length(y2),
  N3 = length(y3),
  N4 = length(y4),
  y1 = y1,
  y2 = y2,
  y3 = y3,
  y4 = y4
)
```

Now that we have the data to pass into our sampling, let's proceed.

```
fit <- mod$sample(
  data = input_data,
  seed = 123,
  chains = 4,
  parallel_chains = 4,
  refresh = 500, # print update every 500 iterations
  iter_warmup = 1000,
  iter_sampling = 5000
)
```

Results

Since we don't care about all the values, we are going to select the specific ones we care about.

```
fit$summary(c("alpha4", "alpha3", "alpha2", "alpha1", "onetwo", "onethree",
              "onefour", "twothree", "twofour", "threefour"))
```


A tibble: 10 x 10

	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk	ess_tail
	<chr>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>
1	alpha4	-3.01	-3.01	0.822	0.799	-4.36	-1.67	1.00	19167.	14643.
2	alpha3	4.02	4.01	0.894	0.857	2.56	5.48	1.00	17767.	13792.
3	alpha2	1.99	2.00	0.884	0.856	0.529	3.43	1.00	19997.	14321.
4	alpha1	-3.00	-3.00	1.03	0.998	-4.69	-1.30	1.00	27457.	15674.
5	onetwo	-4.99	-5.00	1.61	1.54	-7.67	-2.33	1.00	25684.	15976.
6	onethree	-7.02	-7.00	1.62	1.58	-9.66	-4.37	1.00	23064.	15532.
7	onefour	0.0144	0.0129	1.54	1.49	-2.52	2.52	1.00	24007.	16103.
8	twothree	-2.03	-2.02	1.43	1.38	-4.41	0.311	1.00	17113.	12301.
9	twofour	5.01	5.01	1.34	1.30	2.80	7.23	1.00	18883.	13018.
10	threefour	7.03	7.03	1.36	1.32	4.81	9.28	1.00	17591.	13736.


Plot Intervals

Now that we have the draws from the model, we can use them to plot intervals.

```
fit$draws(variables = c("alpha4", "alpha3", "alpha2", "alpha1")) |>
  mcmc_intervals()
```

```
fit$draws(variables = c("mu1", "mu2", "mu3", "mu4", "onetwo", "onethree",
                        "onefour", "twothree", "twofour", "threefour")) |>
  mcmc_intervals()
```



Part VIII

UNIT 8

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.

Part IX

UNIT 9

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.

Part X

UNIT 10

Placeholder

Warning

I have not made it to this point, so this is currently just a placeholder. Please check back regularly as we are updating the files.